

APPLICATIONS



OF DATA SCIENCE

Intro to Web Scrapping

Applications of Data Science - Class Bonus

Giora Simchoni

`gsimchoni@gmail.com` and add `#dsapps` in subject

Stat. and OR Department, TAU

2021-12-31

APPLICATIONS



OF DATA SCIENCE

The Three Rules of Web Scraping

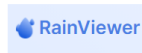
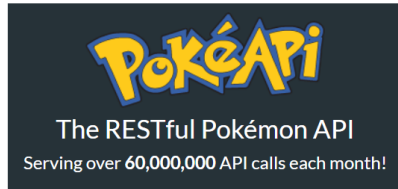
APPLICATIONS



OF DATA SCIENCE

Rule 1: Do you *really* need web scraping?

There are data APIs for just about anything, you know...



R API Packages

Many of them already accessible with a R/Python package...

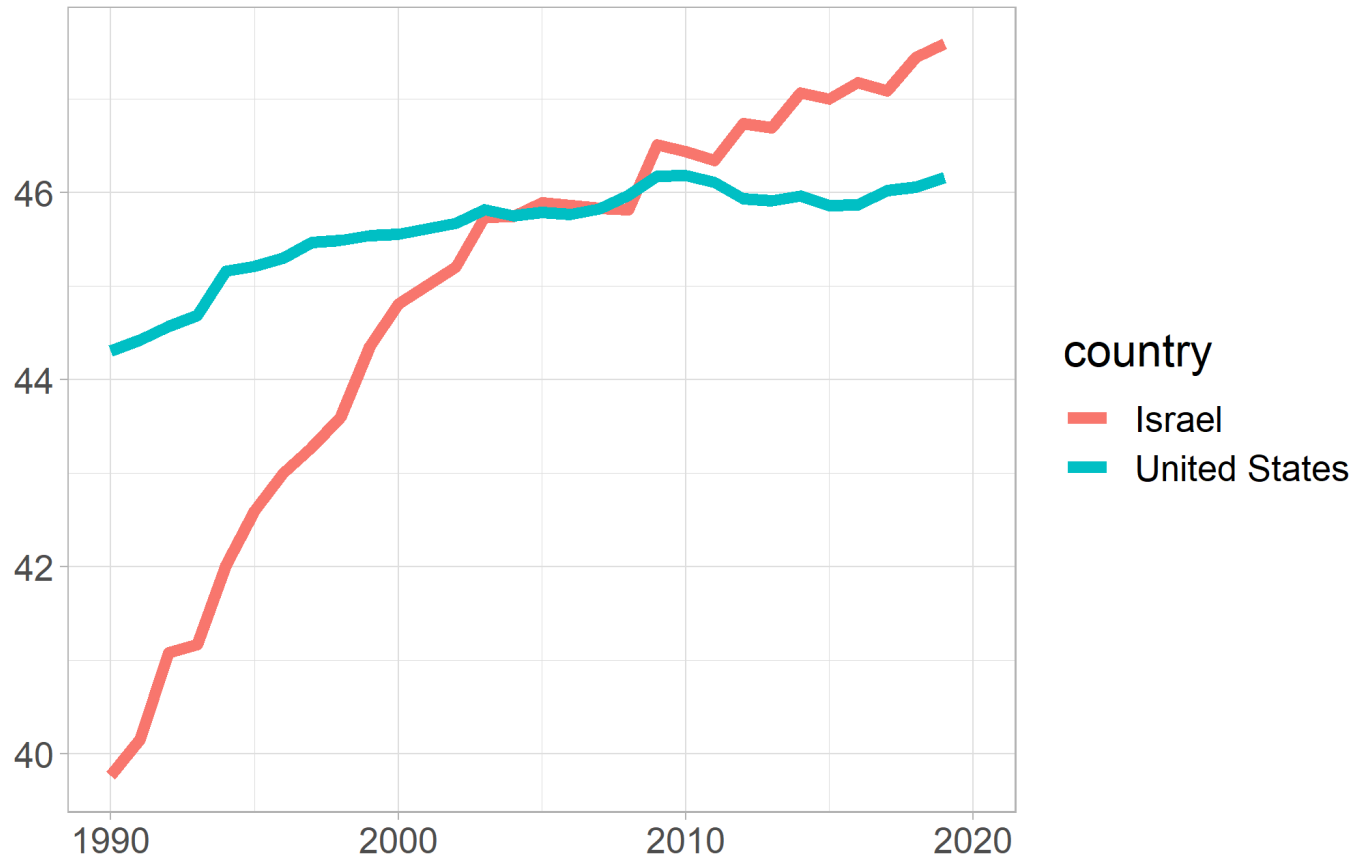
```
library(wbstats)

female_labor <- wb_data(
  indicator = c("women_lab_share" = "SL.TLF.TOTL.FE.ZS"),
  start_date = 1990,
  end_date = 2020
)

female_labor %>%
  filter(country %in% c("Israel", "United States")) %>%
  ggplot(aes(date, women_lab_share, color = country)) +
  geom_line(lwd = 2) +
  labs(title = "Share of women in labor force") +
  theme_light() +
  theme(text = element_text(size=16))
```

From: <https://cfss.uchicago.edu/notes/application-program-interface/>

% of women in labor force



country

- Israel
- United States

The datapasta package

My gift to you.



Rule 2: Learn some HTML first!

HTML is a set (or tree) of *elements*, marked by *HTML tags*:

```
<!DOCTYPE html>
<html>
<head>
<title>My Awesome Webpage</title>
</head>
<body>

<h1>My Superb Heading</h1>
<p>I am going to be a web designer.</p>

</body>
</html>
```

My Superb Heading

I am going to be a web designer.

- First children in the tree: header and body
- View any page's HTML (on chrome) with right-click and "View page source" (or Ctrl + U)

Useful elements and attributes to know

- `<p>` for paragraph `</p>`
- `<h1>` for headings `</h1>`
- `
`, `<hr>` for breaks
- `` for links ``
- `<i>` For bold, italic etc. `</i>`
- ``
- `<p style="color:DodgerBlue;">` for font color `</p>`

HTML Tables

A big thing when it comes to data as you can imagine...

```
<!DOCTYPE html>
<html>
<body>

<h2>Basic HTML Table</h2>

<table style="width:100%">
  <tr>
    <th>Firstname</th>
    <th>Lastname</th>
    <th>Age</th>
  </tr>
  <tr>
    <td>Britney</td>
    <td>Spears</td>
    <td>39</td>
  </tr>
  <tr>
    <td>Christina</td>
    <td>Aguillera</td>
    <td>40</td>
  </tr>
  <tr>
    <td>Beyonce</td>
    <td>Knowles</td>
    <td>39</td>
  </tr>
</table>

</body>
</html>
```

Basic HTML Table

Firstname	Lastname	Age
Britney	Spears	39
Christina	Aguillera	40
Beyonce	Knowles	39

HTML Classes

A class attribute is defined in a style sheet, lets you repeat a style.

```
<!DOCTYPE html>
<html>
<head>
<style>
.city {
  background-color: tomato;
  color: white;
  border: 2px solid black;
  margin: 20px;
  padding: 20px;
}
</style>
</head>
<body>

<div class="city">
<h2>London</h2>
<p>London is the capital of England.</p>
</div>

<div class="city">
<h2>Paris</h2>
<p>Paris is the capital of France.</p>
</div>

<div class="city">
<h2>Tokyo</h2>
<p>Tokyo is the capital of Japan.</p>
</div>

</body>
</html>
```

London

London is the capital of England.

Paris

Paris is the capital of France.

Tokyo

Tokyo is the capital of Japan.

Rule 3: Be polite!

With great power comes great responsibility.

See e.g. the [polite](#) package.



rvest

APPLICATIONS



OF DATA SCIENCE

read_html()

You're now a NLP expert, and you've just developed a SOTA Q&A model. How would you demonstrate your model's performance?

How about [triviaquestionsnow.com](https://www.triviaquestionsnow.com/)?

Let's scrape a few Q&As. Politely.

```
library(rvest)

url <- "https://www.triviaquestionsnow.com/for/sports-trivia"

html_obj <- read_html(url)
```

`read_html()` is usually where you'd start. What did you get?

```
class(html_obj)
```

```
## [1] "xml_document" "xml_node"
```

View page source

With time, you'll become friendly with this weird object. Right now it is better than...

```
1 <!DOCTYPE html>
2 <html lang="en" ng-app="triviaApp">
3 <head>
4   <meta charset="utf-8" />
5   <meta http-equiv="x-ua-compatible" content="ie=edge">
6   <link rel="icon" href="https://www.triviaquestionsnow.com/favicon.ico" type="image/x-icon">
7   <meta name="viewport" content="width=device-width, initial-scale=1.0" />
8   <meta name="description" content="Want to put your sports knowledge to the test? triviaquestionsnow is the right place for you. Work your br>
9   <meta name="keywords" content="Sports Trivia, Sports Trivia questions">
10
11   <link rel="apple-touch-icon" sizes="180x180" href="/apple-touch-icon.png">
12   <link rel="icon" type="image/png" sizes="32x32" href="/favicon-32x32.png">
13   <link rel="icon" type="image/png" sizes="16x16" href="/favicon-16x16.png">
14   <link rel="manifest" href="/site.webmanifest">
15   <link rel="mask-icon" href="/safari-pinned-tab.svg" color="#5bbad5">
16   <meta name="msapplication-TileColor" content="#da532c">
17   <meta name="theme-color" content="#ffffff">
18
19   <link rel="canonical" href="https://www.triviaquestionsnow.com/for/sports-trivia" />
20
21   <meta property="fb:app_id" content="329788460996850" />
22   <meta property="og:url" content="https://www.triviaquestionsnow.com/for/sports-trivia" />
23   <meta property="og:type" content="website" />
24     <meta property="og:title" content="Sports Trivia Questions and Answers" />
25     <meta property="og:description" content="Want to put your sports knowledge to the test? triviaquestionsnow is the right place for>
26     <meta property="og:image" content="https://www.triviaquestionsnow.com/trivia-questions-and-answers.jpg" />
27   <title>Sports Trivia Questions and Answers | TQN</title>
28   <!-- Fonts -->
29   <link rel="stylesheet" href="https://cdnjs.cloudflare.com/ajax/libs/font-awesome/4.5.0/css/font-awesome.min.css" integrity="sha384-XdYbMnZ/Q>
30   <link href='https://fonts.googleapis.com/css?family=Source+Sans+Pro:400,700,300' rel='stylesheet' type='text/css'>
31   <link href="https://cdnjs.cloudflare.com/ajax/libs/foundationicons/3.0.0/foundation-icons.css" rel="stylesheet">
32   <!-- Styles -->
33   <link href="https://www.triviaquestionsnow.com/css/all.css?v=1" rel="stylesheet">
```

html_children() and html_node()

Our tree has two children: head and body

```
html_obj %>% html_children()
```

```
## {xml_nodeset (2)}  
## [1] <head>\n<meta http-equiv="Content-Type" content="text/html; charset=  
## [2] <body>\n      <div class="title-bar" data-responsive-toggle="nav" data
```

Again notice the object returned might not be familiar
("xml_nodeset")

And each of the children has children of its own:

```
html_obj %>% html_node("body") %>% html_children()
```

```
## {xml_nodeset (6)}  
## [1] <div class="title-bar" data-responsive-toggle="nav" data-hide-for="m  
## [2] <div id="nav" class="top-bar">\n      <div class="row">\n  
## [3] <div class="wrap bg-grey-light t-pad-20 b-pad-20">\n      <div  
## [4] <script src="https://www.triviaquestionsnow.com/js/all.js?v=1"></scr
```


html_nodes()

Usually we'd figure out a rule and want a list of all relevant nodes:

```
html_obj %>% html_nodes("img")
```

```
## {xml_nodeset (8)}  
## [1]   
## [2]   
## [3]   
## [4]   
## [5]   
## [6]   
## [7]   
## [8] 
```

```
html_obj %>% html_nodes("a")
```

```
## {xml_nodeset (44)}  
## [1] <a href="/">\n                  
## [2] <a href="https://www.triviaquestionsnow.com" class="no-pad">\n  
## [3] <a href="https://www.triviaquestionsnow.com/easy-trivia-questions">Easy Trivia Questions  
## [4] <a href="https://www.triviaquestionsnow.com/for/sports-trivia">Sports Trivia  
## [5] <a href="https://www.triviaquestionsnow.com/for/music-trivia">Music Trivia
```

html_attr()

Getting a specific attribute from those nodes:

```
html_obj %>% html_nodes("img") %>% html_attr("src")
```

```
## [1] "https://www.triviaquestionsnow.com/img/trivia-questions.png"
## [2] "https://www.triviaquestionsnow.com/img/trivia-questions.png"
## [3] "https://www.triviaquestionsnow.com/img/category/360x130/-category-1"
## [4] "https://www.triviaquestionsnow.com/img/category/360x130/-category-1"
## [5] "https://www.triviaquestionsnow.com/img/category/360x130/-category-1"
## [6] "https://www.triviaquestionsnow.com/img/category/360x130/apologetics"
## [7] "https://www.triviaquestionsnow.com/img/category/360x130/-category-1"
## [8] "https://www.triviaquestionsnow.com/img/category/360x130/-category-2"
```

```
html_obj %>% html_nodes("a") %>% html_attr("href")
```

```
## [1] "/"
## [2] "https://www.triviaquestionsnow.com"
## [3] "https://www.triviaquestionsnow.com/easy-trivia-questions"
## [4] "https://www.triviaquestionsnow.com/for/sports-trivia"
## [5] "https://www.triviaquestionsnow.com/for/music-trivia"
## [6] "https://www.triviaquestionsnow.com/for/math-trivia"
## [7] "https://www.triviaquestionsnow.com/categories"
```

html_text()

Getting the text from whatever set of elements we have:

```
html_obj %>% html_nodes("a") %>% html_text()
```

```
## [1] "\n"
## [2] "\n"
## [3] "Easy Trivia"
## [4] "Sports Trivia"
## [5] "Music Trivia"
## [6] "Math Trivia"
## [7] "Categories"
## [8] "All Trivia"
## [9] "\n"
## [10] "Show Answer"
## [11] "\n"
## [12] "Show Answer"
## [13] "\n"
## [14] "Show Answer"
## [15] "\n"
## [16] "Show Answer"
## [17] "\n"
## [18] "Show Answer"
## [19] "\n"
## [20] "Show Answer"
```

"

"

What male tennis player has won the most Grand S

In what state is the Pro Football Hall of Fame l

Which team won the first Super Bowl ever?\n

Which player from the 1998 NFL Draft is consider

In American football, how many points is a touch

Andre Agassi lost the Wimbledon Championship fin

How to get to those questions? Option 1

Look at the page source, get some identifier yourself (class, ID, link)

Who was the only person to have won a Super Bowl as a player, as an assistant coach and as a head coach? VeryHard

Show Answer

```
<div class="question callout" ng-controller="QuestionController as question">
<input type="hidden" ng-model="question.data.id" ng-init="question.data.id='1443'" />
<span class="float-right light-grey bold 1-cush-10">VeryHard</span>
<h3 class="fs-1 bold">
  <a href="https://www.triviaquestionsnow.com/question/person-to-win-super-bowl-as-player-as-assistant-coach-and-a-head-coach/">Who was the only person to have won a Super Bowl as a player, as an assistant coach and as a head coach?
</h3>
<div class="t-pad-10">
  <a href="#" class="click-to-show bold" ng-click="question.clickShow($event)">Show Answer</a>
  <div class="clearfix"></div>
  <p class="answer hide b-cush-0 t-cush-10" ng-class="{ 'answer': true, 'hide': !(question.showAnswer) }">
    <input type="hidden" name="question.position.1443" value="1" />
    <input type="hidden" ng-model="question.data.id" ng-init="question.data.id='1443'" />
    <input type="hidden" ng-model="question.data.answer" ng-init="question.data.answer='mike-ditka'" />
    <strong>Answer: </strong><em>Mike Ditka</em>
    <br />
    <span class="dark-grey"><!--<strong>More information: </strong-->Ditka won as a player with the Cowboys
  </p>
</div>
</div>
```

```
html_obj %>% html_nodes(".question") %>% .[[1]]
```

```
## {html_node}
## <div class="question callout" ng-controller="QuestionController as quest
## [1] <input type="hidden" ng-model="question.data.id" ng-init="question.d
## [2] <span class="float-right light-grey bold 1-cush-10">Medium</span>
## [3] <h3 class="fs-1 bold">\n          <a href="https://www.triviaquest
## [4] <div class="t-pad-10">\n          <a href="#" class="click-to-show
```

After some trial and error...

```
html_obj %>%  
  html_nodes(".question") %>%  
  html_nodes(".fs-1") %>%  
  html_text() %>%  
  str_trim()
```

```
## [1] "What male tennis player has won the most Grand Slam titles?"  
## [2] "In what state is the Pro Football Hall of Fame located?"  
## [3] "Which team won the first Super Bowl ever?"  
## [4] "Which player from the 1998 NFL Draft is considered by many to be t  
## [5] "In American football, how many points is a touchdown worth?"  
## [6] "Andre Agassi lost the Wimbledon Championship final in 1999. Which  
## [7] "Andre Agassi won his first Olympic gold medal in which year?"  
## [8] "Which male golfer was the winner of 2018 U.S. Open Champion?"  
## [9] "Novak Djokovic won the 2013 Australian Open tournament. Who did he  
## [10] "Up until 2018, who was the only player in NFL history to have rush
```

How to get to those questions? Option 2

[SelectorGadget!](#)



From here it's a function fest!

```
extract_questions_and_answers_from_page <- function(url) {  
  html_obj <- read_html(url)  
  levels <- html_obj %>% html_nodes(".question") %>%  
    html_nodes(".l-cush-10") %>% html_text()  
  questions <- html_obj %>% html_nodes(".question") %>%  
    html_nodes(".fs-1") %>% html_text() %>% str_trim()  
  answers <- html_obj %>% html_nodes(".question") %>%  
    html_nodes(".answer") %>% html_text() %>%  
    str_extract(., "Answer:.*") %>% str_replace("Answer: ", "")  
  tibble(level = levels, question = questions, answer = answers)  
}  
  
extract_questions_and_answers_from_page(url)
```

```
## # A tibble: 10 x 3  
##   level      question                                     answer  
##   <chr>      <chr>                                     <chr>  
## 1 Easy      "The winning team of the Davis Cup is called?"      World  
## 2 Hard      "What was the duration of the longest playoff drough~ 44 year  
## 3 Hard      "In which year did Fred Couples win his first major ~ 1992  
## 4 Hard      "Which player graced the cover of the videogame \"Ma~ Rob Gr  
## 5 Medium    "In 2010, which NBA player posed nude for an issue o~ Amare  
## 6 VeryHard  "Which Super Bowl had the most points ever scored?" Super  
## 7 Easy      "Which country won the 2012 Fed Cup?"               Czech  
## 8 Easy      "Bernhard Langer is a distinguished golfer. Which ge~ German
```

Pagination

```
create_page_url <- function(topic, page_num) {  
  str_c("https://www.triviaquestionsnow.com/for/", topic, "-trivia")  
}  
  
extract_multiple_pages_single_topic <- function(topic, n = 5) {  
  cat(topic, "\n")  
  res <- map_dfr(  
    1:n,  
    function(i) {  
      cat(" ", i)  
      url <- create_page_url(topic, i)  
      extract_questions_and_answers_from_page(url)  
    }  
  )  
  res$topic <- topic  
  cat("\n")  
  res  
}
```



```
extract_multiple_pages_single_topic("sports")
```

```
## # A tibble: 50 x 4
##   level      question                                answer
##   <chr>      <chr>                                <chr>
## 1 VeryHard  "In 1948, which NBA basketball team did the H~ Minneapolis L
## 2 Medium    "The Jacksonville Jaguars and Carolina Panthe~ 1995
## 3 Hard      "An automatic progression by a player to the ~ Bye
## 4 VeryHard  "In 2016, Giants' wide receiver Odell Beckham~ Code Black
## 5 Hard      "Which NBA player broke the record for most p~ Jeremy Lin
## 6 Medium    "Before relocating to Foxborough, Massachuset~ Boston
## 7 Easy      "What is the term for the historic jerseys to~ Throwback Jer
## 8 Medium    "Who served as the starting center of the Gol~ Andrew Bogut
## 9 Easy      "In what year was the 4 minute mile achieved?" 1954 by Roger
## 10 Hard     "Who was the first tennis player to complete ~ Don Budge
## # ... with 40 more rows
```

Magic!

```
topics <- c("sports", "kids", "science", "bible", "food-drink", "r")

df_all <- map_dfr(
  topics,
  extract_multiple_pages_single_topic
)

df_all %>% count(topic)
```

```
## # A tibble: 8 x 2
##   topic          n
##   <chr>        <int>
## 1 bible          50
## 2 food-drink     50
## 3 geography      50
## 4 history        50
## 5 kids           50
## 6 science        50
## 7 sports         50
## 8 video-games    50
```

BeautifulSoup

APPLICATIONS



OF DATA SCIENCE

Almost always start with

```
import requests
from bs4 import BeautifulSoup

html_obj = requests.get('https://en.wikipedia.org/wiki/List_of_The

soup = BeautifulSoup(html_obj.content, 'html.parser')
type(soup)
```

```
## <class 'bs4.BeautifulSoup'>
```

This object has all sorts of attributes and methods:

```
soup.get_text()
soup.prettify()
soup.attrs
soup.children
soup.title
```

find() a tag, find_all()

```
link_objs = soup.find_all('a', href=True)
type(link_objs)
```

```
## <class 'bs4.element.ResultSet'>
```

```
type(link_objs[3])
```

```
## <class 'bs4.element.Tag'>
```

```
link_objs[3].text
```

```
## 'media franchise'
```

```
link_objs[3].attrs
```

```
## {'href': '/wiki/Media_franchise', 'title': 'Media franchise'}
```

See the actual link in the [page](#).

Getting that table

```
table = soup.find('table', attrs={'class': 'wikitable'})
table_body = table.find('tbody')

rows = table_body.find_all('tr')
print(len(rows))
```

```
## 142
```

```
print(rows[0])
```

```
## <tr>
## <th rowspan="2">Installment
## </th>
## <th rowspan="2">Housewives
## </th>
## <th rowspan="2">First season<br/>starred
## </th>
## <th rowspan="2">Last season<br/>starred
## </th>
## <th colspan="3">Number of seasons
## </th></tr>
```

Getting a Housewife name

```
<tr>
<td><span data-sort-value="De La Rosa, Jo&#160;!">Jo De La Rosa</span>
</td>
<td>1
</td>
<td>2
</td>
<td>2
</td>
<td>0
</td>
<td>2
</td></tr>
<tr>
<td><span data-sort-value="Gunvalson, Vicki&#160;!"><a href="/wiki/Vicki_Gunvalson" title="Vicki Gunvalson">Vicki Gunvalson</a></span>
</td>
<td>1
</td>
<td>13
```

```
import re
```

```
print(rows[3].find('span', attrs = {'data-sort-value': re.compile
```

```
## <span data-sort-value="De La Rosa, Jo&nbsp;!">Jo De La Rosa</span>
```

Getting only HWives with Wiki pages

```
housewives_with_links = []
for row in rows:
    housewife = row.find('span',
        attrs = {'data-sort-value': re.compile(r'.*')})
    if housewife is not None:
        link = housewife.find('a')
        if link is not None:
            housewives_with_links.append((housewife.text, link['href']))

import pandas as pd

h_df = pd.DataFrame(housewives_with_links, columns=['name', 'link'])

h_df.head()
```

	name	link
## 0	Vicki Gunvalson	/wiki/Vicki_Gunvalson
## 1	Jeana Keough	/wiki/Jeana_Keough
## 2	Heather Dubrow	/wiki/Heather_Dubrow
## 3	Shannon Storms Beador	/wiki/Shannon_Storms_Beador
## 4	Bethenny Frankel	/wiki/Bethenny_Frankel

(Though if your table is simple, try:)

```
l = pd.read_html(html_obj.text)

l[0].head()
```

```
##      Installment      Housewives ... Number of seasons
##      Installment      Housewives ...      Friend Guest
## 0  Orange County  Kimberly Bryant ...           0.0    3.0
## 1  Orange County    Jo De La Rosa ...           0.0    2.0
## 2  Orange County  Vicki Gunvalson ...           1.0    0.0
## 3  Orange County    Jeana Keough  ...           1.0    4.0
## 4  Orange County    Lauri Peterson ...           1.0    1.0
##
## [5 rows x 7 columns]
```

Following HWives Links

```
def get_housewife_img_ref(housewife_link):
    html_obj = requests.get('https://en.wikipedia.org' + housewife_link)
    soup = BeautifulSoup(html_obj.content, 'html.parser')
    infobox = soup.find('table', attrs = {'class': 'vcard'})
    if infobox is not None:
        img_obj = infobox.find('img', src=True)
        if img_obj is not None:
            return img_obj['src']
    return None

h_df['img_ref'] = h_df['link'].apply(get_housewife_img_ref)
h_df.dropna(inplace=True)

h_df.head()
```

```
##              name    ...
## 0      Vicki Gunvalson  ... //upload.wikimedia.org/wikipedia/common
## 1      Luann de Lesseps  ... //upload.wikimedia.org/wikipedia/common
## 2      Bethenny Frankel  ... //upload.wikimedia.org/wikipedia/common
## 3  Kelly Killoren Bensimon  ... //upload.wikimedia.org/wikipedia/common
## 4      Carole Radziwill  ... //upload.wikimedia.org/wikipedia/common
##
## [5 rows x 3 columns]
```

Downloading HWives Images

```
def make_img_filename(hf_name):  
    return 'data/housewives/' + hf_name.lower().strip(',.-').replace  
  
def download_hw_img(hf_name, hf_img_ref):  
    img_file = make_img_filename(hf_name)  
    img_data = requests.get('http:' + hf_img_ref).content  
    with open(img_file, 'wb') as handler:  
        handler.write(img_data)  
  
h_df.apply(lambda row: download_hw_img(row['name'], row['img_ref']  
axis=1)
```